

Query lifecycle and the planner

IMPROVING QUERY PERFORMANCE IN POSTGRESQL

SQL

Amy McCarty
Instructor

Basic query lifecycle

	System	Front end steps	Back end processes
1	Parser	Send query to database	Checks syntax. Translates SQL into more computer friendly syntax based on system stored rules.
2	Planner & Optimizer	Assess and optimize query tasks	Uses database stats to create query plan. Calculates costs and chooses the best plan.
3	Executor	Return query results	Follows the query plan to execute the query.

Query planner and optimizer

Responsive to SQL structure changes

- Generates plan trees
 - Nodes corresponding to steps
 - Visualize with EXPLAIN
- Estimate cost of each tree
 - Statistics from pg_tables
 - Time based optimization

¹ Plan tree: <https://www.postgresql.org/docs/current/querytree.html>

Statistics from pg_tables

```
SELECT * FROM pg_class  
WHERE relname = 'mytable'
```

```
-- sample of output columns  
| relname | relhasindex |
```

```
SELECT * FROM pg_stats  
WHERE tablename = 'mytable'
```

```
-- sample of output columns  
null_frac | avg_width | n_distinct |
```

- Column indexes
- Count null values
- Column width
- Distinct values

EXPLAIN

- Window into query plan
- Steps and cost **estimates**
 - Does not run query
- Sequential scan of cheeses table
- Cost and size estimates

```
EXPLAIN
```

```
SELECT * FROM cheeses
```

```
Seq Scan on cheeses
```

```
(cost=0.00..10.50 rows=5725 width=296)
```

EXPLAIN: Scan

- Query plan step
- Returns rows

Seq Scan on cheeses (cost=0.00..10.50
rows=5725 width=296)

- **Seq Scan** : scan of all the rows in table

EXPLAIN: Cost

- Dimensionless
- Compare structures with same output
 - Should **not** compare queries with different output

Seq Scan on cheeses (**cost=0.00..10.50**
rows=5725 width=296)

- **0.00..** : start up time
- **..10.50** : total time
- total time = start up + run time

EXPLAIN: Size

- Size estimates

Seq Scan on cheeses (cost=0.00..10.50
rows=5725 width=296)

- **rows** : rows query needs to examine to run
- **width** : byte width of rows

EXPLAIN with a WHERE clause

EXPLAIN

```
SELECT * FROM cheeses WHERE species IN ('goat', 'sheep')
```

```
Seq Scan on cheeses (cost=0.00..378.90 rows=3 width=118)  
-> Filter: (species = ANY ('{"goat","sheep"}'::text[]))
```

- From bottom to top
 - Step 1: Filter
 - Step 2: Sequential scan
- WHERE clause
 - Decrease rows to scan and increases total cost

EXPLAIN with an index

EXPLAIN

```
SELECT * FROM cheeses WHERE species IN ('goat', 'sheep') -- index on species column
```

```
Bitmap Index Scan using species_idx on cheeses (cost=0.29..12.66 rows=3 width=118)  
  Index Cond: (species = ANY ('{"goat","sheep"}'::text[]))
```

- Step 1: Bitmap Index Scan
 - Index Cond explains the scan step
- INDEX
 - Start up cost increased from 0
 - Overall cost decreased from 379

Let's practice!

IMPROVING QUERY PERFORMANCE IN POSTGRESQL

A deeper dive into EXPLAIN

IMPROVING QUERY PERFORMANCE IN POSTGRESQL



Amy McCarty
Instructor

EXPLAIN optional parameters

VERBOSE

- Columns for each plan node
- Shows table schema and aliases

ANALYZE

- Runs the query
- Actual run times in milliseconds

VERBOSE

```
EXPLAIN VERBOSE
```

```
SELECT * FROM cheeses
```

```
Seq Scan on dairy.cheeses (cost=0.00..10.50 rows=5725 width=296)
```

```
Output: name, species, type, age
```

ANALYZE

EXPLAIN ANALYZE

SELECT * FROM cheeses

```
Seq Scan on cheeses (cost=0.00..10.50 rows=5725 width=296) (actual  
time=0.007..1.087 rows=11992 loops=1)  
Planning Time: 0.059 ms  
Execution Time: 1.538 ms
```

- Most useful to minimize run time

Query plan - aggregations

```
EXPLAIN ANALYZE
```

```
SELECT type, AVG(age) AS avg_age
```

```
FROM cheeses
```

```
GROUP BY type -- hard or soft cheese
```

```
HashAggregate (cost=314.88..317.38 rows=200 width=40)(actual time = 4.973..4.975  
              rows=2 loops=1)
```

```
  Group Key: type
```

```
    -> Seq Scan on cheeses (cost=0.00..286.25 rows=5725 width=10)(actual time =  
                                0.016..2.546 rows = 11992 loops=1)
```

```
Planning Time: 12.891 ms
```

```
Execution Time: 5.074 ms
```


Query plan - sort

```
EXPLAIN ANALYZE
```

```
SELECT name, age
```

```
FROM cheeses
```

```
ORDER BY age DESC
```

```
Sort  (cost=1161.37..1191.35 rows=11992 width=20)(actual time = 4.281..5.331
      rows=11992 loops=1)
  Sort Key: age DESC
  Sort Method: quicksort Memory: 1216kB
-> Seq Scan on cheeses  (cost=0.00..348.92 rows=11992 width=20)(actual time =
      0.0007..1.799 rows = 11992 loops=1)

Planning Time: 0.131 ms
Execution Time: 5.870 ms
```

EXPLAIN ANALYZE

```
SELECT name, age FROM cheeses  
INNER JOIN animals ON cheeses.species = animals.species
```

```
Hash Join  (cost=182.97..4339.35 rows=335776 width=145)(actual time=2.755..138.418  
          rows=335776 loops=1)  
  Hash Cond: (cheeses.species = animals.species)  
    -> Seq Scan on cheeses  (cost=0.00..348.92 rows=11992 width=118) (actual  
          time=0.010..2.271 rows=11992 loops=1)  
    -> Hash  (cost=106.32..106.32 rows=6132 width=27) (actual time=2.725..2.725 rows=6132  
          loops=1)  
          Buckets: 8192  Batches: 1  Memory Usage: 439kB  
          -> Seq Scan on animals  (cost=0.00..106.32 rows=6132 width=27) (actual  
                time=0.009..1.008 rows=6132 loops=1)  
  
Planning Time: 0.379 ms  
Execution Time: 161.918 ms
```

Let's practice!

IMPROVING QUERY PERFORMANCE IN POSTGRESQL

Query structure and query execution

IMPROVING QUERY PERFORMANCE IN POSTGRESQL

SQL

Amy McCarty
Instructor

Subqueries and joins

-- SUBQUERY

```
SELECT COUNT(athlete_id)
FROM athletes
WHERE country IN
  (SELECT country FROM climate
   WHERE temp_annual > 22)
```

-- JOIN

```
SELECT COUNT(athlete_id)
FROM athletes a
INNER JOIN climate c
  ON a.country = c.country
  AND c.temp_annual > 22
```

Query plan

```
Aggregate ()  
  -> Hash Join ()  
    Hash Cond: (athletes.country = climate.country)  
    -> Seq Scan on athletes ()  
    -> Hash ()  
      -> Seq Scan on climate ()  
        Filter: (temp_annual > '22'::numeric)
```

Common table expressions and temporary tables

```
-- CTE
WITH celsius AS
(
  SELECT country
  FROM climate
  WHERE temp_annual > 22 -- Celsius
)
SELECT count(athlete_id)
FROM athletes a
INNER JOIN celsius c
  ON a.country = c.country
```

```
-- TEMP TABLE
CREATE TEMPORARY TABLE celsius AS
  SELECT country
  FROM climate
  WHERE temp_annual > 22; -- Celsius

SELECT count(athlete_id)
FROM athletes a
INNER JOIN celsius c
  ON a.country = c.country
```

Query plan

```
Aggregate ()  
CTE celsius  
  -> Seq Scan on climate ()  
      Filter: (temp_annual > '22'::numeric)  
-> Hash Join ()  
    Hash Cond: (a.country_code = c.country_code)  
    -> Seq Scan on athletes a ()  
    -> Hash ()  
        -> CTE Scan on celsius c ()
```


Limiting the data

```
SELECT country_code  
  , COUNT(athlete_id) as athletes  
FROM athletes  
WHERE year in (2014, 2010) -- Indexed column  
GROUP BY country_code
```

Limiting the data

```
SELECT country_code
, COUNT(athlete_id) as athletes
FROM athletes
WHERE year in (2014, 2010) -- Indexed column
GROUP BY country_code
```

No Index	Index
Planning Time: 3.370 ms	Planning Time: 0.163 ms
Execution Time: 0.143 ms	Execution Time: 0.062 ms

Aggregations - different granularities

```
SELECT r.country
      , COUNT(a.athlete_id) as athletes
FROM regions r -- country level
INNER JOIN athletes a -- athletes level
  ON r.country = a.country
GROUP BY r.country
```

- Execution Time : 0.267 ms

Aggregations - changing the granularity

```
WITH olympians AS ( -- country level
  SELECT country
    , COUNT(athlete_id) as athletes
  FROM athletes -- athletes level
  GROUP BY country
)
SELECT country, athletes
FROM regions r -- country level
INNER JOIN olympians o
  ON r.country = o.country
```

	Execution Time
Join 1st	0.267 ms
Aggregate 1st	0.192 ms

Let's practice!

IMPROVING QUERY PERFORMANCE IN POSTGRESQL


Congratulations

IMPROVING QUERY PERFORMANCE IN POSTGRESQL



Amy McCarty
Instructor

Chapter 1: ways to combine data

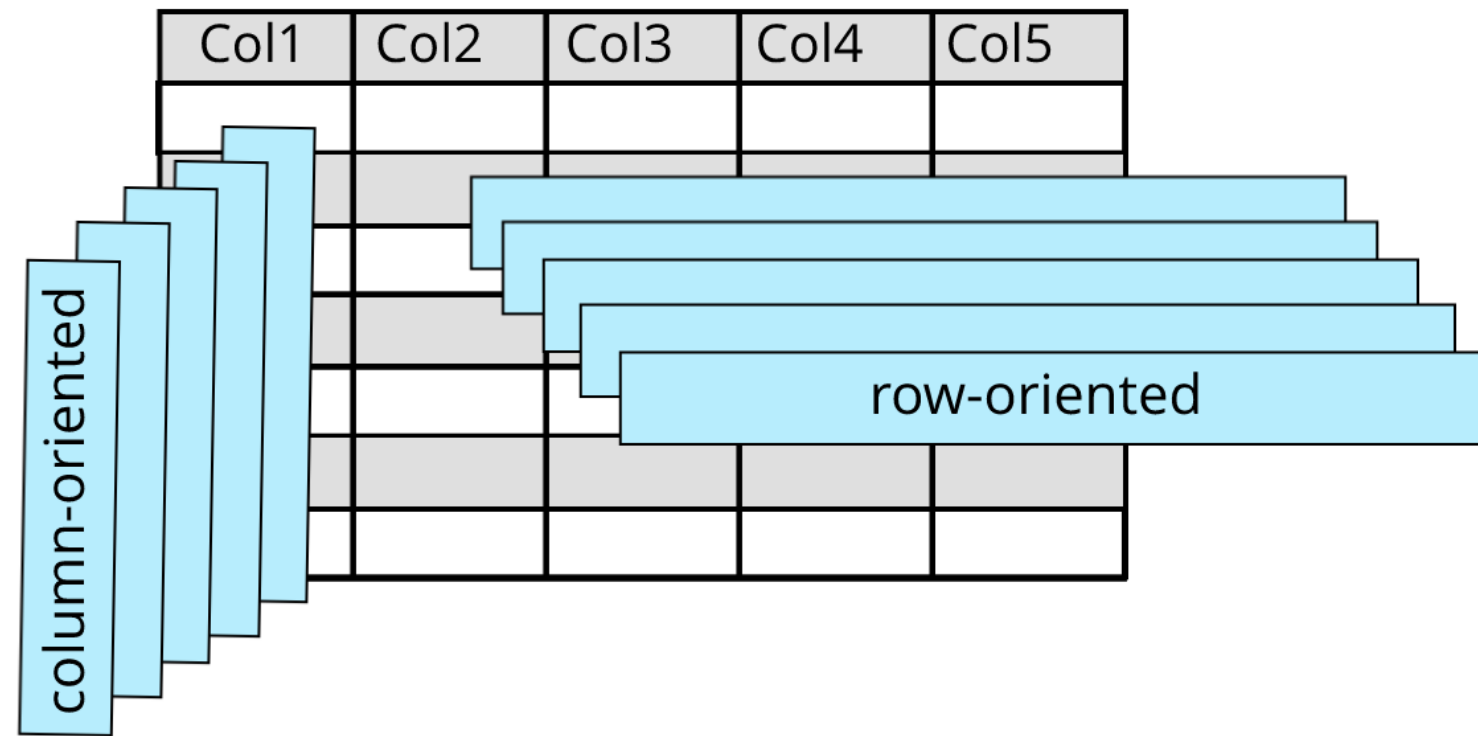


Chapter 2: limiting the results

- Order of operations
- Filtering in the WHERE
- Filtering with an INNER JOIN
- Joining different data granularities

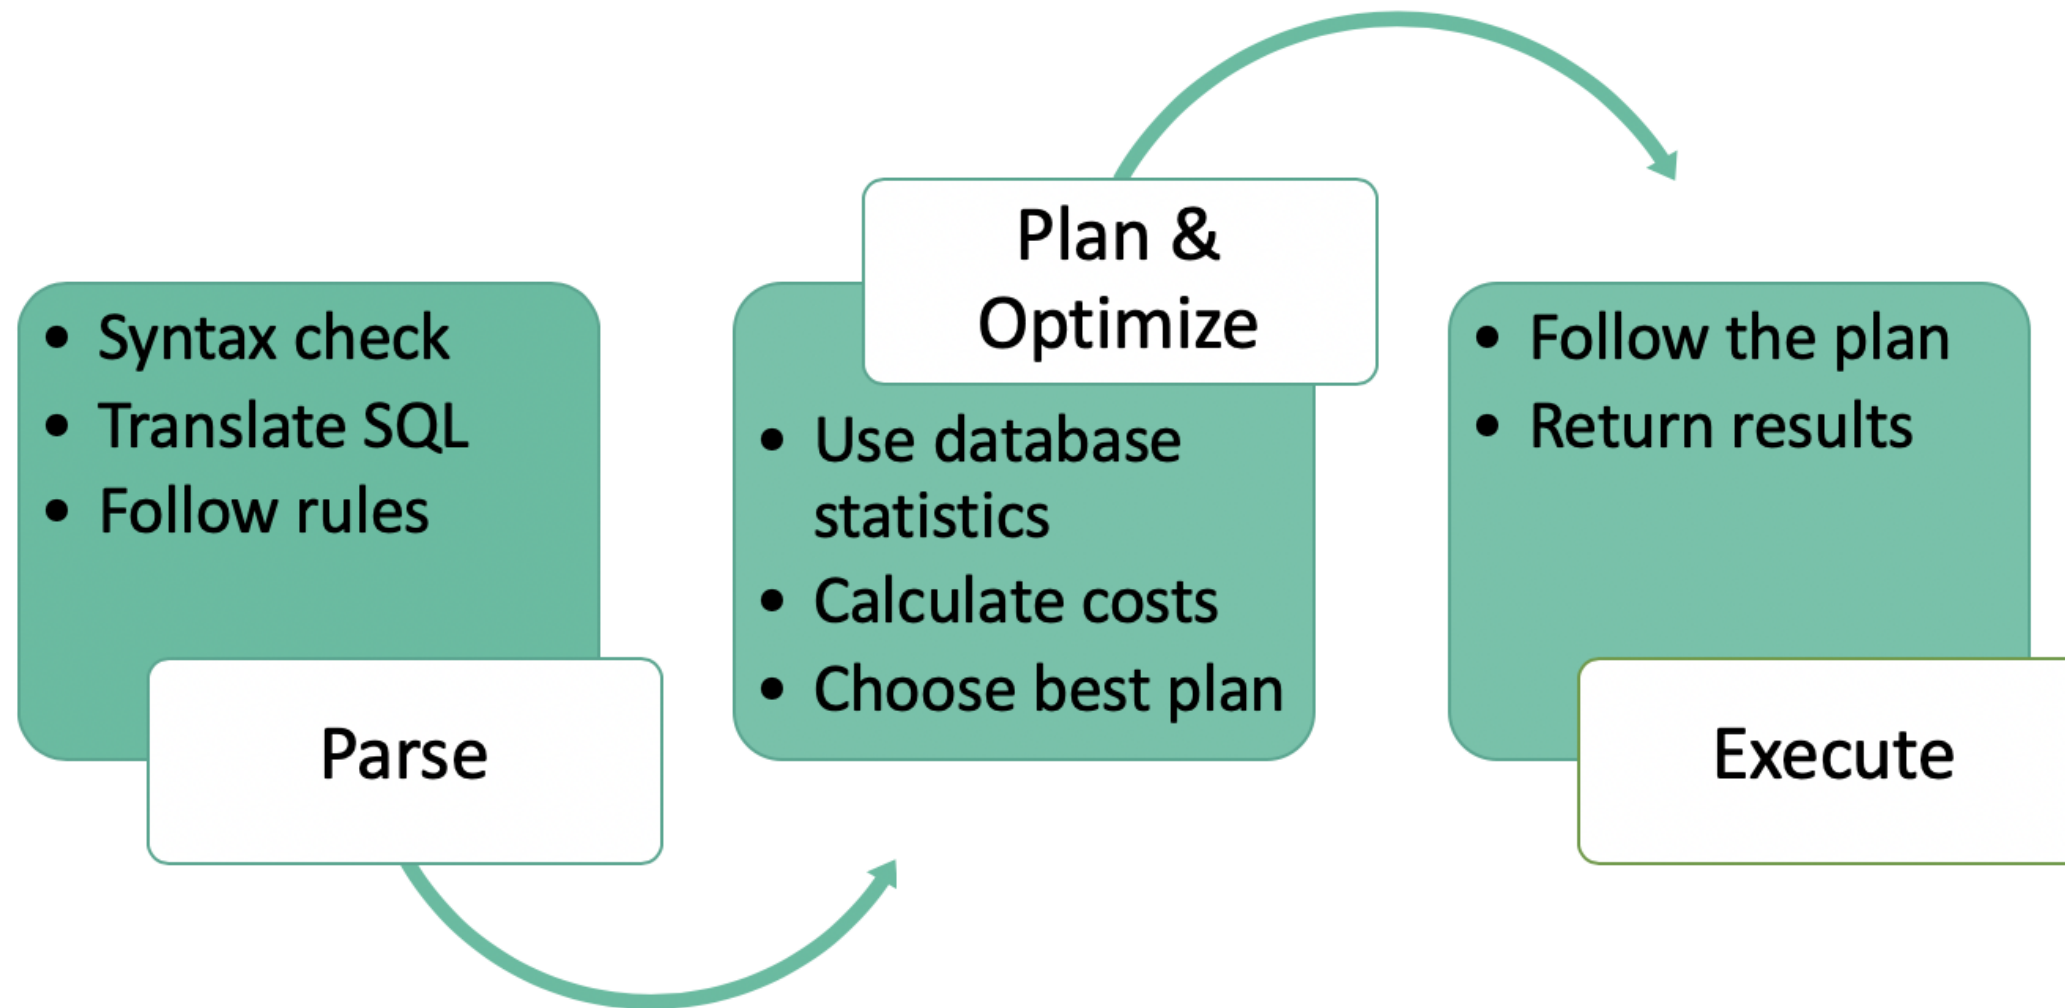
Chapter 3: learning the database

- pg_tables and information_schema
- Tables and views



- Indexes and partitions

Chapter 4: using the query planner



- EXPLAIN

Thank you!

IMPROVING QUERY PERFORMANCE IN POSTGRESQL